

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier

from sklearn.metrics import accuracy_score, classification_report

import warnings
warnings.filterwarnings("ignore")
import plotly.express as px
import plotly.io as pio
```

```
df=pd.read_csv("/content/bank (1).csv")
```

```
df
```

	age	job	marital	education	default	balance	housing	loan	contact
0	59	0	1	1	0	2343	1	0	
1	56	0	1	1	0	45	0	0	
2	41	9	1	1	0	1270	1	0	
3	55	7	1	1	0	2476	1	0	
4	54	0	1	2	0	184	0	0	
...
11157	33	1	2	0	0	1	1	0	
11158	39	7	1	1	0	733	0	0	
11159	32	9	2	1	0	29	0	0	
11160	43	9	1	1	0	0	0	1	

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11162 entries, 0 to 11161
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         11162 non-null  int64
1   job         11162 non-null  int64
2   marital     11162 non-null  int64
3   education   11162 non-null  int64
4   default     11162 non-null  int64
5   balance     11162 non-null  int64
6   housing     11162 non-null  int64
7   loan        11162 non-null  int64
8   contact     11162 non-null  int64
9   day         11162 non-null  int64
10  month       11162 non-null  int64
11  duration    11162 non-null  int64
12  campaign    11162 non-null  int64
13  pdays       11162 non-null  int64
14  previous    11162 non-null  int64
15  poutcome    11162 non-null  int64
16  deposit     11162 non-null  int64
dtypes: int64(17)
memory usage: 1.4 MB
```

```
df.isnull().sum()
```

```
age      0
job      0
```

```
marital      0
education    0
default      0
balance      0
housing      0
loan         0
contact      0
day          0
month        0
duration     0
campaign     0
pdays      0
previous     0
poutcome     0
deposit      0
dtype: int64
```

df.describe()

	age	job	marital	education	default
count	11162.000000	11162.000000	11162.000000	11162.000000	11162.000000
mean	41.231948	4.487905	1.199337	1.285164	0.015051
std	11.913369	3.225132	0.625552	0.749478	0.121761
min	18.000000	0.000000	0.000000	0.000000	0.000000
25%	32.000000	1.000000	1.000000	1.000000	0.000000
50%	39.000000	4.000000	1.000000	1.000000	0.000000
75%	49.000000	7.000000	2.000000	2.000000	0.000000
max	95.000000	11.000000	2.000000	3.000000	1.000000

df.corr()

	age	job	marital	education	default	balance
age	1.000000	-0.031603	-0.442782	-0.126018	-0.011425	0.112300
job	-0.031603	1.000000	0.078314	0.147046	-0.007066	0.028736
marital	-0.442782	0.078314	1.000000	0.125845	-0.014691	-0.002138
education	-0.126018	0.147046	0.125845	1.000000	-0.010709	0.051728
default	-0.011425	-0.007066	-0.014691	-0.010709	1.000000	-0.060954
balance	0.112300	0.028736	-0.002138	0.051728	-0.060954	1.000000
housing	-0.168700	-0.136965	-0.036345	-0.109168	0.011076	-0.077092
loan	-0.031418	-0.067092	-0.062029	-0.073154	0.076434	-0.084589
contact	0.027762	-0.087915	-0.060456	-0.132540	0.035709	-0.027295
day	-0.000762	0.026589	-0.003642	0.016759	0.017342	0.010467
month	-0.026130	-0.076011	-0.004070	-0.055868	0.000950	0.007264
duration	0.000189	0.002432	0.006781	-0.019122	-0.009760	0.022436
campaign	-0.005278	0.003104	-0.030794	-0.005327	0.030975	-0.013894
pdays	0.002774	-0.003385	0.031200	0.025165	-0.036282	0.017411
previous	0.020169	0.012665	0.031281	0.022427	-0.035273	0.030805

df.drop(["job","marital","education","default","contact","campaign"],axis=1,inplace=True)

df.head()

	age	balance	housing	loan	day	month	duration	pdays	previous	poutcome
0	59	2343	1	0	5	8	1042	-1	0	
1	56	45	0	0	5	8	1467	-1	0	
2	41	1270	1	0	5	8	1389	-1	0	
3	55	2476	1	0	5	8	579	-1	0	

df.keys

<bound	method	NDFrame.keys	of	age	balance	housing	loan	day	month	duration	pdays	previous	\
0	59	2343	1	0	5	8	1042	-1	0				
1	56	45	0	0	5	8	1467	-1	0				
2	41	1270	1	0	5	8	1389	-1	0				
3	55	2476	1	0	5	8	579	-1	0				
4	54	184	0	0	5	8	673	-1	0				
...				
11157	33	1	1	0	20	0	257	-1	0				
11158	39	733	0	0	16	6	83	-1	0				
11159	32	29	0	0	19	1	156	-1	0				
11160	43	0	0	1	8	8	9	172	5				
11161	34	0	0	0	9	5	628	-1	0				
	poutcome	deposit											
0	3	1											
1	3	1											
2	3	1											
3	3	1											
4	3	1											
...											
11157	3	0											
11158	3	0											
11159	3	0											
11160	0	0											
11161	3	0											

[11162 rows x 11 columns]>

sns.pairplot(df)




```
df.isin(df)
```

```
age balance housing loan day month duration pdays previous
```

```
df.interpolate()
```

	age	balance	housing	loan	day	month	duration	pdays	previous
0	59	2343	1	0	5	8	1042	-1	0
1	56	45	0	0	5	8	1467	-1	0
2	41	1270	1	0	5	8	1389	-1	0
3	55	2476	1	0	5	8	579	-1	0
4	54	184	0	0	5	8	673	-1	0
...
11157	33	1	1	0	20	0	257	-1	0
11158	39	733	0	0	16	6	83	-1	0
11159	32	29	0	0	19	1	156	-1	0
11160	43	0	0	1	8	8	9	172	5

```
plt.figure(figsize= (15,6))
sns.set_style("darkgrid")
sns.heatmap(df.corr(),annot= True)
plt.show()
```




```
ax= px.histogram(df,x= "loan", template= "plotly_dark",color= "previous",title='Age distribution')
ax.show()
```

Age distribution



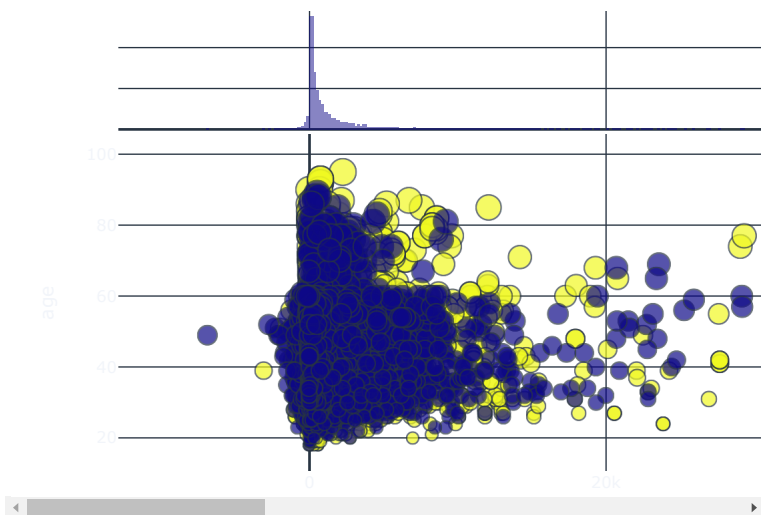
```
ax= px.pie(df, names= "housing",template= "plotly_dark",title= "housing loan distrbusion",hole= 0.5)
ax.show()
```

housing loan distrbusion




```
ax= px.scatter(df,x= "balance",y= "age",marginal_x='histogram', marginal_y='histogram',size="age", size_max=20,
    template= "plotly_dark",color= "deposit",title="age and diposite correlation")
ax.show()
```

age and diposite correlation




```
ax= px.scatter(df,x= "age",y= "housing",marginal_x='histogram', marginal_y='histogram',size="age", size_max=20,
    template= "plotly_dark",color= "previous",title="age and housingcorrelation")
ax.show()
```


age and housingcorrelation




```
ax = px.scatter_3d(df, x="age", y="loan", z="deposit",template= "plotly_dark",color="month")
ax.show()
```



```
ax= px.parallel_coordinates(df, color="loan",template= "plotly_dark")
ax.show()
```



```
ax = px.box(df,template= "plotly_dark")
ax.show()
```



```
#splittig x and y for predictoions
```

```
x=df.iloc[:,0:-1]
x
```

	age	balance	housing	loan	day	month	duration	pdays	previous
0	59	2343	1	0	5	8	1042	-1	0
1	56	45	0	0	5	8	1467	-1	0
2	41	1270	1	0	5	8	1389	-1	0
3	55	2476	1	0	5	8	579	-1	0
4	54	184	0	0	5	8	673	-1	0
...
11157	33	1	1	0	20	0	257	-1	0
11158	39	733	0	0	16	6	83	-1	0
11159	32	29	0	0	19	1	156	-1	0
11160	43	0	0	1	8	8	9	172	5


```
y=df["deposit"]
```

```
y
```

```
0      1
1      1
2      1
3      1
4      1
```

```
..
```

```
11157    0
11158    0
11159    0
11160    0
11161    0
```

```
Name: deposit, Length: 11162, dtype: int64
```

```
xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.20,random_state=1)
```

```
lreg=LogisticRegression()
```

```
knn=KNeighborsClassifier()
```

```
svm=SVC()
```

```
dt=DecisionTreeClassifier()
```

```
def mymodel(model):
```

```
    model.fit(xtrain,ytrain)
```

```
    ypred=model.predict(xtest)
```

```
    ac=accuracy_score(ytest,ypred)
```

```
    cr=classification_report(ytest,ypred)
```

```
    print(f"accuracy-: {ac} \n classification report\n {cr}")
```

```
mymodel(lreg)
```

```
accuracy-: 0.7617554858934169
```

```
classification report
```

	precision	recall	f1-score	support
0	0.76	0.79	0.78	1165
1	0.76	0.73	0.75	1068
accuracy			0.76	2233
macro avg	0.76	0.76	0.76	2233
weighted avg	0.76	0.76	0.76	2233

```
mymodel(knn)
```

```
accuracy-: 0.7523510971786834
```

```
classification report
```

	precision	recall	f1-score	support
0	0.75	0.79	0.77	1165
1	0.75	0.72	0.73	1068
accuracy			0.75	2233
macro avg	0.75	0.75	0.75	2233
weighted avg	0.75	0.75	0.75	2233

```
mymodel(svm)
```

```
accuracy-: 0.7393640841916704
```

```
classification report
```

	precision	recall	f1-score	support
0	0.71	0.84	0.77	1165
1	0.78	0.63	0.70	1068
accuracy			0.74	2233
macro avg	0.75	0.73	0.73	2233
weighted avg	0.75	0.74	0.74	2233

```
mymodel(dt)
```

```
accuracy-: 0.7756381549484997
```

```
classification report
```

	precision	recall	f1-score	support
0	0.79	0.78	0.78	1165
1	0.76	0.78	0.77	1068
accuracy			0.78	2233
macro avg	0.78	0.78	0.78	2233
weighted avg	0.78	0.78	0.78	2233

```
#checking overfitted on dt
dt.score(xtrain,ytrain)
```

```
1.0
```

```
dt1=DecisionTreeClassifier(max_depth=5)
mymodel(dt1)
```

accuracy-: 0.7877295118674429				
classification report				
	precision	recall	f1-score	support
0	0.77	0.84	0.81	1165
1	0.81	0.73	0.77	1068
accuracy			0.79	2233
macro avg	0.79	0.79	0.79	2233
weighted avg	0.79	0.79	0.79	2233

```
for i in range(1,50):
    dt2=DecisionTreeClassifier(max_depth=i)
    dt2.fit(xtrain,ytrain)
    ypred=dt2.predict(xtest)
    print(f"{i}= {accuracy_score(ytest,ypred)}")
```

```
1= 0.7098074339453649
2= 0.7098074339453649
3= 0.7666815942678011
4= 0.7787729511867443
5= 0.7890729959695477
6= 0.7957904164800716
7= 0.8007165248544559
8= 0.8011643528884909
9= 0.8034034930586654
10= 0.8101209135691895
11= 0.8025078369905956
12= 0.7984773846842812
13= 0.800268696820421
14= 0.7980295566502463
15= 0.7948947604120018
16= 0.780564263322884
17= 0.7819077474249888
18= 0.780564263322884
19= 0.7769816390506046
20= 0.7760859829825347
21= 0.7693685624720108
22= 0.77384684281236
23= 0.77384684281236
24= 0.7747424988804299
25= 0.7747424988804299
26= 0.7729511867442902
27= 0.7711598746081505
28= 0.7671294223018361
29= 0.7720555306762203
30= 0.7792207792207793
31= 0.7778772951186744
32= 0.7756381549484997
33= 0.7765338110165696
34= 0.774294670846395
35= 0.7733990147783252
36= 0.77384684281236
37= 0.7689207344379758
38= 0.7693685624720108
39= 0.77384684281236
40= 0.768025078369906
41= 0.767577250335871
42= 0.7733990147783252
43= 0.77384684281236
44= 0.7725033587102552
```

```

45= 0.7733990147783252
46= 0.774294670846395
47= 0.77384684281236
48= 0.768025078369906
49= 0.7733990147783252

```

```

#best value of max_depth=10
dt3=DecisionTreeClassifier(max_depth=11)
mymodel(dt3)

```

```

accuracy-: 0.8034034930586654
classification report

```

	precision	recall	f1-score	support
0	0.81	0.81	0.81	1165
1	0.79	0.80	0.80	1068
accuracy			0.80	2233
macro avg	0.80	0.80	0.80	2233
weighted avg	0.80	0.80	0.80	2233

```

dt4=DecisionTreeClassifier(min_samples_leaf=10)
mymodel(dt4)

```

```

accuracy-: 0.8047469771607703
classification report

```

	precision	recall	f1-score	support
0	0.81	0.82	0.81	1165
1	0.80	0.79	0.79	1068
accuracy			0.80	2233
macro avg	0.80	0.80	0.80	2233
weighted avg	0.80	0.80	0.80	2233

```

for i in range(1,75):
dt5=DecisionTreeClassifier(min_samples_leaf=i)
dt5.fit(xtrain,ytrain)
ypred=dt5.predict(xtest)
print(f"{i}= {accuracy_score(ytest,ypred)}")

```

```

15= 0.8110165696372593
16= 0.8056426332288401
17= 0.8087774294670846
18= 0.8092252575011196
19= 0.8141513658755039
20= 0.8145991939095387
21= 0.8101209135691895
22= 0.8101209135691895
23= 0.8101209135691895
24= 0.8119122257053292
25= 0.8083296014330497
26= 0.8101209135691895
27= 0.8087774294670846
28= 0.8096730855351545
29= 0.8096730855351545
30= 0.8078817733990148
31= 0.8096730855351545
32= 0.8078817733990148
33= 0.8105687416032243
34= 0.8119122257053292
35= 0.8101209135691895
36= 0.8105687416032243
37= 0.8114643976712942

```

```
53= 0.8150470219435737
54= 0.8150470219435737
55= 0.8154948499776086
56= 0.8159426780116436
57= 0.8168383340797134
58= 0.8168383340797134
59= 0.8168383340797134
60= 0.819525302283923
61= 0.819525302283923
62= 0.819525302283923
63= 0.819525302283923
64= 0.819525302283923
65= 0.8154948499776086
66= 0.8159426780116436
67= 0.8092252575011196
68= 0.8092252575011196
69= 0.8132557098074339
70= 0.8132557098074339
71= 0.812807881773399
72= 0.812807881773399
73= 0.8105687416032243

dt6=DecisionTreeClassifier(min_samples_leaf=41)
mymodel(dt6)

accuracy-: 0.8177339901477833
classification report
      precision    recall  f1-score   support

      0       0.85       0.79       0.82       1165
      1       0.79       0.84       0.82       1068

   accuracy       0.82
  macro avg       0.82
 weighted avg       0.82

dt7=DecisionTreeClassifier(max_depth=11,min_samples_leaf=41)
mymodel(dt7)

accuracy-: 0.8159426780116436
classification report
      precision    recall  f1-score   support

      0       0.84       0.79       0.82       1165
      1       0.79       0.84       0.81       1068

   accuracy       0.82
  macro avg       0.82
 weighted avg       0.82

#gini
dt8=DecisionTreeClassifier(criterion="gini",min_samples_leaf=41)
mymodel(dt8)

accuracy-: 0.8177339901477833
classification report
      precision    recall  f1-score   support

      0       0.85       0.79       0.82       1165
      1       0.79       0.84       0.82       1068

   accuracy       0.82
  macro avg       0.82
 weighted avg       0.82

dt9=DecisionTreeClassifier(criterion="gini",max_depth=10)
mymodel(dt9)

accuracy-: 0.8105687416032243
classification report
      precision    recall  f1-score   support

      0       0.83       0.80       0.82       1165
      1       0.79       0.82       0.80       1068

   accuracy       0.81
  macro avg       0.81
```

```
weighted avg      0.81      0.81      0.81      2233
```

```
dt10=DecisionTreeClassifier(criterion="entropy",max_depth=10)
mymodel(dt10)
```

```
accuracy-: 0.8150470219435737
classification report
      precision    recall  f1-score   support

      0       0.81      0.84      0.83      1165
      1       0.82      0.79      0.80      1068

 accuracy
macro avg      0.82      0.81      0.81      2233
weighted avg    0.82      0.82      0.81      2233
```

```
for i in range(1,50):
    dt11=DecisionTreeClassifier(criterion="entropy",max_depth=i)
    dt11.fit(xtrain,ytrain)
    ypred=dt11.predict(xtest)
    print(f"{i}= {accuracy_score(ytest,ypred)}")
```

```
1= 0.7098074339453649
2= 0.7098074339453649
3= 0.7532467532467533
4= 0.7711598746081505
5= 0.7823555754590238
6= 0.8007165248544559
7= 0.8096730855351545
8= 0.8150470219435737
9= 0.819525302283923
10= 0.8168383340797134
11= 0.8096730855351545
12= 0.8078817733990148
13= 0.7957904164800716
14= 0.7895208240035826
15= 0.7850425436632333
16= 0.7877295118674429
17= 0.7895208240035826
18= 0.7913121361397224
19= 0.793999104343932
20= 0.7877295118674429
21= 0.7850425436632333
22= 0.7850425436632333
23= 0.7877295118674429
24= 0.7854903716972682
25= 0.781012091356919
26= 0.786833855799373
27= 0.7836990595611285
28= 0.7886251679355127
29= 0.7863860277653381
30= 0.780564263322884
31= 0.7859381997313032
32= 0.7850425436632333
33= 0.786833855799373
34= 0.7836990595611285
35= 0.7890729959695477
36= 0.7774294670846394
37= 0.7859381997313032
38= 0.7823555754590238
39= 0.7854903716972682
40= 0.7828034034930587
41= 0.7823555754590238
42= 0.781012091356919
43= 0.7859381997313032
44= 0.7854903716972682
45= 0.7881773399014779
46= 0.7845947156291984
47= 0.7823555754590238
48= 0.786833855799373
49= 0.7819077474249888
```

```
#best value of max_depth when criterion=entropy
dt12=DecisionTreeClassifier(criterion="entropy",max_depth=10)
mymodel(dt12)
```

```
accuracy-: 0.8172861621137483
classification report
```

	precision	recall	f1-score	support
0	0.82	0.84	0.83	1165
1	0.82	0.79	0.81	1068
accuracy			0.82	2233
macro avg	0.82	0.82	0.82	2233
weighted avg	0.82	0.82	0.82	2233

```
dt13=DecisionTreeClassifier(criterion="entropy",min_samples_leaf=10)
mymodel(dt13)
```

accuracy-: 0.8029556650246306				
classification report				
	precision	recall	f1-score	support
0	0.81	0.81	0.81	1165
1	0.79	0.80	0.79	1068
accuracy			0.80	2233
macro avg	0.80	0.80	0.80	2233
weighted avg	0.80	0.80	0.80	2233

```
for i in range(1,50):
    dt14=DecisionTreeClassifier(criterion="entropy",min_samples_leaf=i)
    dt14.fit(xtrain,ytrain)
    ypred=dt14.predict(xtest)
    print(f"{i}= {accuracy_score(ytest,ypred)}")
```

```
1= 0.7881773399014779
2= 0.7707120465741155
3= 0.7845947156291984
4= 0.7845947156291984
5= 0.7890729959695477
6= 0.793551276309897
7= 0.799820868786386
8= 0.8029556650246306
9= 0.8092252575011196
10= 0.8038513210927004
11= 0.8056426332288401
12= 0.8029556650246306
13= 0.800268696820421
14= 0.8007165248544559
15= 0.802060089565607
16= 0.7984773846842812
17= 0.8034034930586654
18= 0.8016121809225257
19= 0.8083296014330497
20= 0.8110165696372593
21= 0.8096730855351545
22= 0.8132557098074339
23= 0.8101209135691895
24= 0.8083296014330497
25= 0.8141513658755039
26= 0.8150470219435737
27= 0.8172861621137483
28= 0.8163905060456784
29= 0.8163905060456784
30= 0.8163905060456784
31= 0.8145991939095387
32= 0.812807881773399
33= 0.8101209135691895
34= 0.8096730855351545
35= 0.8105687416032243
36= 0.8096730855351545
37= 0.8074339453649798
38= 0.8038513210927004
39= 0.80653828929691
40= 0.8092252575011196
41= 0.8110165696372593
42= 0.8141513658755039
43= 0.8101209135691895
```

```
# best value of min_samples_leaf=15 when criterion=entropy
dt15=DecisionTreeClassifier(criterion="entropy",min_samples_leaf=15)
mymodel(dt15)
```

```

accuracy-: 0.8020600089565607
classification report
              precision    recall  f1-score   support

         0       0.81      0.80      0.81      1165
         1       0.79      0.80      0.79      1068

    accuracy          0.80
   macro avg          0.80
weighted avg          0.80

```

**FINAL MODEL

```

#gini
dt16=DecisionTreeClassifier(max_depth=10,min_samples_leaf=41)
mymodel(dt16)

```

```

accuracy-: 0.8159426780116436
classification report
              precision    recall  f1-score   support

         0       0.84      0.79      0.82      1165
         1       0.79      0.84      0.81      1068

    accuracy          0.82
   macro avg          0.82
weighted avg          0.82

```

```

#entropy
dt17=DecisionTreeClassifier(criterion="entropy",min_samples_leaf=15)
mymodel(dt17)

```

```

accuracy-: 0.8020600089565607
classification report
              precision    recall  f1-score   support

         0       0.81      0.80      0.81      1165
         1       0.79      0.80      0.79      1068

    accuracy          0.80
   macro avg          0.80
weighted avg          0.80

```

```
# import Random Forest classifier
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
# instantiate the classifier
```

```
rfc = RandomForestClassifier(random_state=0)
```

```
# fit the model
```

```
rfc.fit(xtrain, ytrain)
```

```
# Predict the Test set results
```

```
ypred = rfc.predict(xtest)
```

```
# Check accuracy score
```

```
from sklearn.metrics import accuracy_score
```